

project proposal

1. Set up

Packages

```
# general use and cleaning
library(tidyverse)
library(here)
library(janitor)
library(snakecase)

# wrap axis labels
library(scales)

# read .xlsx files
library(readxl)

# visualize missing data
library(naniar)

# arrange plots
library(patchwork)

library(dplyr)      # data manipulation
library(broom)      # tidy statistical output
library(corrplot)
```

Data

```

# taxonomic information
taxon_list <- read_csv(here("data", "taxon_list.csv"))

# aquatic invertebrates
aq_ins <- read_xlsx(here("data", "Aquatic Sampling Data-2026-03-10.xlsx"),
  sheet = "Aquatic Insects")

# read in site information
sites <- read_xlsx(here("data", "Aquatic Sampling Data-2026-03-10.xlsx"),
  sheet = "sites")

# read in water quality information
water_quality <- read_xlsx(here("data", "Aquatic Sampling Data-2026-03-10.xlsx"),
  sheet = "Water Quality",
  na = c("", "n/a", "N/A", "over", "173+"))

```

2. Cleaning

sites object

```

# creating new object from sites
ncos_sites <- sites |>
  # filter to only include NCOS sites sampled four times a year
  filter(sector == "NCOS" & sampling_frequency == "Quarterly")

```

Cleaning taxon_list

```

# creating new clean object from taxon_list
taxon_list_clean <- taxon_list |>
  # converting taxon name to snake case (no spaces or capital letters,
  # only underscores)
  mutate(verbatim_name = to_any_case(verbatim_name, case = "snake"))

```

Cleaning water_quality

```
# creating new clean object from water quality
water_quality_clean <- water_quality |>
  # clean column names
  clean_names() |>
  # filtering join: only include sites in ncos_sites
  semi_join(ncos_sites, by = "site") |>
  # create new column to join with later data frames
  unite("date_site", date, site, remove = TRUE) |>
  # group by date_site column
  group_by(date_site) |>
  # calculate median pH, salinity, DO
  summarize(
    med_ph = median(dissolved_oxygen_mg_l, na.rm = TRUE),
    med_sal = median(salinity_ppt, na.rm = TRUE),
    med_do = median(dissolved_oxygen_mg_l, na.rm = TRUE)) |>
  # ungroup data frame
  ungroup() |>
  # filter out salinity outliers
  filter(date_site != "2022-08-12_NVBR")
```

Cleaning aq_ins

```
# creating new clean object from aquatic inverts
aq_ins_clean <- aq_ins |>
  # clean column names
  clean_names() |>
  # filtering join: only include sites occurring in ncos_sites
  semi_join(ncos_sites, by = c("site")) |>
  # renaming columns
  rename(date = date_on_vial) |>
  # select columns of interest
  select(
    date, sample_type, site,
    ostracod,
    copepod,
    hemiptera_corixidae_boatman,
    cladocera,
    nematode,
    diptera_ceratopogonidae,
    annelida_oligochaete,
```

```

    diptera_chironomid,
    annelida_polychaete) |>
# filter to only include "filtered beaker" samples
filter(sample_type %in% c("FB 250", "FB250")) |>
# convert the data frame to long format
pivot_longer(cols = ostracod:annelida_polychaete,
              names_to = "taxon",
              values_to = "abundance") |>
# group by date, site, taxon
group_by(date, site, taxon) |>
# calculate average abundance per liter (sum abundance divided by 7.5 liters)
summarize(ave_lit = sum(abundance, na.rm = TRUE)/7.5) |>
# ungroup data frame
ungroup() |>
# create a new column called `date_site` to join with water quality
unite("date_site", date, site, remove = FALSE) |>
# join with water quality
left_join(water_quality_clean, by = c("date_site")) |>
# join with taxon list
left_join(taxon_list_clean, by = c("taxon" = "verbatim_name")) |>
# add full names of sites
mutate(site_full = case_when(
  site == "NVBR" ~ "North of Venoco Bridge",
  site == "NEC" ~ "South of Venoco Bridge",
  site == "NMC" ~ "Main Channel",
  site == "NPB" ~ "South of Phelps Bridge",
  site == "NPB1" ~ "North of Phelps Bridge",
  site == "NPB2" ~ "Phelps Road",
  site == "NWP" ~ "West Pond",
  site == "NDC" ~ "Devereux Creek"
)) |>
# setting factor levels for site
mutate(site_full = fct_reorder(site_full, med_do, .fun = "median", .na_rm = TRUE),
       site_full = fct_rev(site_full))

```

Focus Taxa

```

# taxa to include in all figures
focus_taxa <- c("Ostracod", "Copepod", "Cladocera",
                 "Hemiptera Corixidae Boatman", "Nematode")

```

Figure 1: Invertebrate abundance as a function of DO

```
# create object DO_taxa from aq_ins_clean
DO_taxa <- aq_ins_clean |>
  # remove rows that are missing median salinity values
  filter(!is.na(med_do)) |>
  # create column that calculates log-transformed abundance
  mutate(log_ave_lit = log(ave_lit + 1)) |>
  # mutates to title-case taxon names
  mutate(taxon = str_replace_all(taxon, "_", " "), taxon = str_to_title(taxon)) |>
  # only include focus taxa
  filter(taxon %in% focus_taxa) |>
  # order taxon levels by maximum average abundance per liter
  mutate(taxon = fct_reorder(taxon, ave_lit, .fun = max, .desc = TRUE)) |>
  # keep only the columns needed for the figure
  select(date, site, site_full, taxon, ave_lit, log_ave_lit, med_do)

# creates median point location for each taxon
DO_taxa_medians <- DO_taxa |>
  # groups data by taxon so each species gets one median point
  group_by(taxon) |>
  # calculates median DO and median log abundance for each taxon
  summarise(
    med_do_point = median(med_do, na.rm = TRUE),
    med_log_ave_lit = median(log_ave_lit, na.rm = TRUE),
    .groups = "drop"
  )

DO_taxa_plot <- ggplot(data = DO_taxa, # assigns plot to object do_taxa_plot
  # uses DO_taxa data frame
  mapping = aes(x = med_do, # x-axis: med_do
    y = log_ave_lit, # y-axis: log_ave_lit
    color = taxon)) + # color by taxon

# first layer: points show individual abundance measurements
geom_point(shape = 16, # shape is circle
  size = 2, # size is 2
  alpha = 0.7) + # opacity is 0.7

# second layer: larger point shows the median location for each taxon
geom_point(data = DO_taxa_medians,
  mapping = aes(x = med_do_point,
```


```

                                y = med_log_ave_lit),
    shape = 18,
    size = 5) +

# separate panels for each taxon with free x- and y-axis scales
facet_wrap(~ taxon, scales = "free") +
# use non-default colors for points
scale_color_viridis_d() +
labs(x = "Median DO (mg/L)", # relabel x-axis
     y = "log + 1 transformed average abundance/L", # relabel y-axis
     # creates descriptive title
     title = "Invertebrate abundance as a function of DO") +

# uses non-default theme
theme_minimal() +
# remove redundant legend
theme(legend.position = "none")
DO_taxa_plot

```



shows that individual taxa have noisy, weak, or inconsistent relationships with DO.

Figure 2: Variance of DO across sights

```
# create object aq_ins_clean_fig2 from aq_ins_clean
aq_ins_clean_fig2 <- aq_ins_clean |>
  # convert taxon to title case
  mutate(taxon = str_replace_all(taxon, "_", " "),
         taxon = str_to_title(taxon)) |>
  # filter to include focus taxa
  filter(taxon %in% focus_taxa)

# base layer
```

```

DO <- ggplot(data = aq_ins_clean_fig2,
  # x-axis: site
  # y-axis: mean salinity
  mapping = aes(x = site_full,
    y = med_do)) +
# first layer: boxplot
geom_boxplot(width = 0.35,
  outlier.shape = NA) +
# second layer: salinity measurements at each survey
geom_jitter(height = 0,
  width = 0.12,
  shape = 21,
  alpha = 0.6) +
# wrapping x-axis text
scale_x_discrete(labels = label_wrap(width = 15)) +
# second layer: points to represent means
stat_summary(fun = median,
  color = "red",
  size = 1) +
# relabelling x- and y-axes
labs(x = "Site",
  y = "Median DO (mh/L)",
  title = "Site DO levels") +
# different theme
theme_bw()
# show plot
DO

```





Figure 3: Aquatic Invert. Abundance Across Sites

```
# creating new object from clean aquatic inverts
abundance_by_site <- aq_ins_clean |>
  # log + 1 transform mean abundance per liter
  mutate(log_ave_lit = log(ave_lit + 1)) |>
  # convert taxon to title case
  mutate(taxon = str_replace_all(taxon, "_", " "),
         taxon = str_to_title(taxon)) |>
  # filter to include taxa of interest
  filter(taxon %in% focus_taxa)
abundance <- ggplot(data = abundance_by_site,
  # x-axis: site
```

```

# y-axis: mean abundance/L
# color by taxon
mapping = aes(x = site_full,
              y = log_ave_lit,
              color = taxon)) +
# first layer: jitter points to represent observations
geom_jitter(width = 0.15,
            height = 0,
            alpha = 0.7) +
# add median point for each site within each taxon
stat_summary(fun = median,
            geom = "point",
            shape = 18,
            size = 3,
            color = "black") +
# facet by taxon
facet_wrap(~ taxon) +
# wrapping x-axis text
scale_x_discrete(labels = label_wrap(15)) +
# non-default colors
scale_color_viridis_d() +
# relabelling x- and y-axis, adding title
labs(x = "Site",
     y = "log + 1 transformed average abundance/L",
     title = "Aquatic Invert. Abundance Across Sites") +
theme_bw() +
# remove legend
theme(legend.position = 'none',
      # fix x-axis for readability
      axis.text.x = element_text(angle = 45, hjust = 1))

```

```

# base layer

```

```

abundance <- ggplot(data = abundance_by_site,
                   # x-axis: site
                   # y-axis: mean abundance/L
                   # color by taxon
                   mapping = aes(x = site_full,
                                y = log_ave_lit,
                                color = taxon)) +
# first layer: jitter points to represent observations
geom_jitter(width = 0.15,

```

```

        height = 0,
        alpha = 0.7) +
# facet by taxon
facet_wrap(~ taxon) +
# wrapping x-axis text
scale_x_discrete(labels = label_wrap(15)) +
# non-default colors
scale_color_viridis_d() +
# relabelling x- and y-axis, adding title
labs(x = "Site",
      y = "log + 1 transformed average abundance/L",
      title = "Aquatic Invert. Abundance Across Sites") +
theme_bw() +
# remove legend
theme(legend.position = 'none',
      # fix x-axis for readability
      axis.text.x = element_text(angle = 45, hjust = 1))
abundance

```




Figure 4: Abundance over time

```
# create object for plotting total abundance over time
abundance_time <- aq_ins_clean |>
  # group by sampling date and taxon
  group_by(date, taxon) |>
  # total abundance across all sites for each date and taxon
  summarize(total_ave_lit = sum(ave_lit, na.rm = TRUE)) |>
  # ungroup data frame
  ungroup() |>
  # create log-transformed total abundance column
  mutate(log_total_ave_lit = log(total_ave_lit + 1)) |>
  # make taxon names easier to read
```

```


mutate(taxon = str_replace_all(taxon, "_", " "),
       taxon = str_to_title(taxon)) |>
# order taxa by maximum total abundance
mutate(taxon = fct_reorder(taxon, total_ave_lit, .fun = max, .desc = TRUE)) |>
# include main taxon groups of focus
filter(taxon %in% focus_taxa)

# base layer
abundance_time_plot <- ggplot(data = abundance_time,
                             # x-axis: log_total_ave_lit
                             # fill and color by taxon
                             mapping = aes(x = log_total_ave_lit,
                                           fill = taxon,
                                           color = taxon)) +
# scaled density makes y-axis consistent across facets
geom_density(aes(y = after_stat(scaled)),
             alpha = 0.45,
             linewidth = 0.8,
             adjust = 0.8) +
facet_wrap(~ taxon) + # facet by taxon
scale_fill_viridis_d() + # non-default colors
scale_color_viridis_d() + # non-default colors
labs(x = "log + 1 transformed total average abundance/L across sites",
     y = "Scaled density",
     title = "Distribution of total invertebrate abundance over time") +
# use clean theme
theme_bw() +
# remove redundant legend
theme(legend.position = "none",
      axis.text.x = element_text(angle = 45,
                                  hjust = 1))

# show plot
abundance_time_plot

```

Distribution of total invertebrate abundance over time



Correlation analysis

```
# create wide dataset (one column per taxon)
taxa_wide <- DO_taxa |>
  select(date, site, taxon, log_ave_lit, med_do) |>
  pivot_wider(
    names_from = taxon,
    values_from = log_ave_lit
  )

# calculate correlation of each taxon with dissolved oxygen
# calculate Spearman correlation of each taxon with dissolved oxygen
cor_do <- taxa_wide |>
```

```

select(-date, -site) |>
summarize(across(-med_do, ~ cor(.x,
                                med_do,
                                method = "spearman",
                                use = "complete.obs")))) |>

pivot_longer(cols = everything(),
              names_to = "taxon",
              values_to = "correlation")


ggplot(cor_do, aes(x = taxon, y = "Dissolved oxygen", fill = correlation)) +
  geom_tile(color = "grey70") +
  geom_point(aes(size = abs(correlation)), shape = 21, fill = "white") +

  # move labels ABOVE circles
  geom_text(aes(label = round(correlation, 2)),
            nudge_y = 0.15, # <-- this moves text up
            size = 4) +

  scale_fill_gradient2(
    low = "red",
    mid = "white",
    high = "blue",
    midpoint = 0,
    limits = c(-1, 1),
    name = "r"
  ) +
  scale_size(range = c(2, 10), guide = "none") +
  labs(
    x = "Taxon",
    y = NULL,
    title = "Spearman correlation between dissolved oxygen and invertebrate abundance"
  ) +
  theme_minimal() +
  theme(
    panel.grid = element_blank(),
    axis.text.x = element_text(angle = 45, hjust = 1)
  )

```

Spearman correlation between dissolved oxygen and ir



```
# Spearman correlation between DO and abundance for each taxon
spearman_by_taxon <- DO_taxa |>
  # keep only species of interest
  filter(taxon %in% focus_taxa) |>
  # group by taxon so each species gets its own test
  group_by(taxon) |>
  # run Spearman correlation test for each taxon
  summarise(
    spearman_rho = cor(log_ave_lit,
                      med_do,
                      method = "spearman",
                      use = "complete.obs"),
    p_value = cor.test(log_ave_lit,
                      med_do,
                      method = "spearman",
                      exact = FALSE)$p.value,
    n = sum(complete.cases(log_ave_lit, med_do)),
    .groups = "drop"
  )

# show results
spearman_by_taxon
```



```
# A tibble: 5 x 4
```

	taxon	spearman_rho	p_value	n
	<fct>	<dbl>	<dbl>	<int>
1	Copepod	-0.335	0.0529	34
2	Ostracod	-0.128	0.470	34
3	Cladocera	-0.342	0.0475	34
4	Hemiptera Corixidae Boatman	-0.0869	0.625	34
5	Nematode	0.0105	0.953	34

Summed abundance across focus taxa vs DO

```
focus_taxa <- c("Cladocera", "Copepod",
                 "Hemiptera Corixidae Boatman",
                 "Nematode", "Ostracod")


total_abundance <- DO_taxa |>
  filter(taxon %in% focus_taxa) |>
  group_by(date, site, med_do) |>
  summarize(
    total_abundance = sum(ave_lit, na.rm = TRUE),
    log_total_abundance = log10(total_abundance + 1),
    .groups = "drop"
  )
ggplot(total_abundance,
       aes(x = med_do,
           y = log_total_abundance)) +

  geom_point(color = "magenta3",
             size = 3,
             alpha = 0.7) +
  geom_smooth(method = "lm",
             se = TRUE,
             color = "purple4",
             fill = "cyan2") +

  labs(
    x = "Dissolved oxygen (mg/L)",
    y = "Log total invertebrate abundance",
    title = "Relationship between dissolved oxygen and total invertebrate abundance"
  ) +
```

```
theme_classic()
```

Relationship between dissolved oxygen and total invertebrate abundance



When taxa are combined, a clearer negative relationship emerges.

```
cor.test(  
  total_abundance$log_total_abundance,  
  total_abundance$med_do,  
  method = "spearman",  
  exact = FALSE)
```

Spearman's rank correlation rho

```
data: total_abundance$log_total_abundance and total_abundance$med_do  
S = 8526.3, p-value = 0.08182  
alternative hypothesis: true rho is not equal to 0  
sample estimates:  
rho  
-0.3027201
```

Binned DO levels

```
total_abundance <- total_abundance |>
  mutate(
    do_category = case_when(
      med_do <= quantile(med_do, 0.33, na.rm = TRUE) ~ "Low DO",
      med_do <= quantile(med_do, 0.66, na.rm = TRUE) ~ "Medium DO",
      TRUE ~ "High DO"
    )
  )
total_abundance$do_category <- factor(
  total_abundance$do_category,
  levels = c("Low DO", "Medium DO", "High DO")
)

ggplot(total_abundance,
  aes(x = do_category,
      y = log_total_abundance,
      fill = do_category)) +


  geom_boxplot(alpha = 0.7,
    outlier.color = "black",
    outlier.size = 2) +

  # optional: add jittered points to show raw data
  geom_jitter(width = 0.15,
    alpha = 0.5,
    size = 2,
    color = "black") +

  labs(
    x = "Dissolved oxygen level",
    y = "Log total invertebrate abundance",
    title = "Invertebrate abundance across dissolved oxygen levels"
  ) +

  scale_fill_manual(values = c(
    "Low DO" = "red3",
    "Medium DO" = "grey60",
    "High DO" = "blue3"
  )) +
```

```
theme_classic() +
theme(legend.position = "none")
```



Community abundance varies continuously rather than discretely across DO levels.

```
# Pearson correlation for continuous DO and total community abundance
cor.test(total_abundance$log_total_abundance,
         total_abundance$med_do,
         method = "pearson")
```

Pearson's product-moment correlation

```
data: total_abundance$log_total_abundance and total_abundance$med_do
t = -2.2882, df = 32, p-value = 0.02887
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 -0.63289820 -0.04217141
sample estimates:
cor
-0.3749894
```

```
# r = -0.3749894, p = 0.02887
# there is a significant negative correlation between abundance and continuous DO

# Kruskal-Wallis test comparing abundance across DO categories
kruskal.test(log_total_abundance ~ do_category,
             data = total_abundance)
```

Kruskal-Wallis rank sum test

```
data: log_total_abundance by do_category
Kruskal-Wallis chi-squared = 3.2212, df = 2, p-value = 0.1998
```

```
# Kruskal-Wallis chi-squared = 3.2212, df = 2, p-value = 0.1998
# there is not a significant difference between abundance and DO categories
```

Composition plot across DO bins

```
composition_plot <- DO_taxa |>
  mutate(
    do_category = case_when(
      med_do <= quantile(med_do, 0.33, na.rm = TRUE) ~ "Low DO",
      med_do <= quantile(med_do, 0.66, na.rm = TRUE) ~ "Medium DO",
      TRUE ~ "High DO"
    ),

    # reorder factor levels
    do_category = factor(do_category,
                        levels = c("Low DO",
                                   "Medium DO",
                                   "High DO"))
  ) |>
  group_by(do_category, taxon) |>
  summarize(
    total_abundance = sum(ave_lit, na.rm = TRUE),
    .groups = "drop"
  ) |>
  ggplot(aes(x = do_category,
             y = total_abundance,
```

```

    fill = taxon)) +

geom_col(position = "fill") +


scale_y_continuous(labels = scales::percent_format()) +

labs(
  x = "Dissolved oxygen category",
  y = "Relative community composition",
  fill = "Taxon",
  title = "Invertebrate community composition shifts across dissolved oxygen levels"
) +

theme_classic()

composition_plot

```




Medium DO sites show a more even distribution of taxa and greater relative abundance of Cladocera and Ostracods, while low and high DO conditions are dominated primarily by Copepods.

Time-series plot (all species)

```
# species/taxa to include in the total abundance figure
species_interest <- c(
  "Ostracod",
  "Copepod",
  "Cladocera",
  "Hemiptera Corixidae Boatman",
  "Nematode"
)

# create total abundance over time
total_abundance_time <- DO_taxa |>
  # keeps only the taxa of interest
  filter(taxon %in% species_interest) |>
  # makes sure date is treated as a date
  mutate(date = as.Date(date)) |>
  # groups by date only, so abundance is summed across sites and taxa
  group_by(date) |>
  # sums average abundance per liter across all taxa of interest
  summarise(
    total_abundance = sum(ave_lit, na.rm = TRUE),
    .groups = "drop"
  )

# plot total abundance over time
ggplot(total_abundance_time, aes(x = date, y = total_abundance)) +
  # adds a line connecting total abundance values over time
  geom_line(linewidth = 1) +
  # adds points for each sampling date
  geom_point(size = 2) +
  # formats y-axis labels with commas
  scale_y_continuous(labels = scales::comma) +
  # adds clear labels
  labs(
    title = "Total Abundance of Species of Interest Over Time",
    x = "Sampling Date",
    y = "Total Abundance"
  ) +
  # uses a clean theme
  theme_minimal()
```



Time-series area plot

```
# species/taxa to include in the area plot
species_interest <- c(
  "Ostracod",
  "Copepod",
  "Cladocera",
  "Hemiptera Corixidae Boatman",
  "Nematode"
)

# create abundance over time by taxon
taxon_abundance_time <- DO_taxa |>
  # keeps only the taxa of interest
  filter(taxon %in% species_interest) |>
  # makes sure date is treated as a date
  mutate(date = as.Date(date)) |>
  # groups by date and taxon so each species stays separate
  group_by(date, taxon) |>
  # sums across sites, but not across taxa
```




```

summarise(
  ave_lit = sum(ave_lit, na.rm = TRUE),
  .groups = "drop"
)

# create time series area plot by taxon
ggplot(taxon_abundance_time, aes(x = date, y = ave_lit, fill = taxon)) +
  # creates stacked area plot with a different color for each taxon
  geom_area(alpha = 0.8) +
  # formats y-axis labels with commas
  scale_y_continuous(labels = scales::comma) +
  # adds clear labels
  labs(
    title = "Abundance of Species of Interest Over Time",
    x = "Sampling Date",
    y = "Average Abundance per Liter",
    fill = "Taxon"
  ) +
  # uses a clean theme
  theme_minimal()

```



DO Bins Across Sites

```
# create new object to store medians of median do levels
median_fig2 <- aq_ins_clean_fig2 |>
  group_by(site_full) |>
  summarize(median_do = median(med_do, na.rm = TRUE))

# assign quartiles for medians for each site
median_fig2$quartile <- cut(median_fig2$median_do,
                           breaks = quantile(median_fig2$median_do,
                                              probs = seq(0, 1, 0.25)),
                           include.lowest = TRUE,
                           labels = c("Q1", "Q2", "Q3", "Q4"))


# boundaries of quartiles
q <- quantile(median_fig2$median_do, probs = c(0.25, 0.5, 0.75))

# distribution plot of DO levels across sites
ggplot(median_fig2, aes(x = median_do)) +
  # density curve
  geom_density(fill = "lightblue",
              alpha = 0.4,
              color = "blue",
              linewidth = 1.2) +
  # vertical quartile lines
  geom_vline(xintercept = q,
            linetype = "dashed",
            color = "gray40") +
  # site labels
  geom_text(aes(y = 0.01,
               label = str_wrap(site_full, width = 10),
               color = quartile),
            hjust = 0.5,
            size = 3) +
  # remove space between bottom of plot and y-axis, set y-axis limits
  scale_y_continuous(expand = c(0,0),
                    limits = c(0, 0.15)) +
  # axes labels and title
  labs(x = "Median Dissolved Oxygen (mg/L)",
       y = "Density",
       title = "Distribution of Median Dissolved Oxygen Across Sites",
```

```

    color = "Quartile") +
  # theme
  theme_bw()

```



BINS Low DO: Phelps Road and North of Phelps Bridge Moderate DO: South of Venoco Bridge, South of Phelps Bridge, Devereux Creek High DO: North of Venoco Bridge, Main Channel

Invertebrate Abundance in DO Bins

```

aq_ins_clean_bins <- aq_ins_clean |>
  mutate(bins = case_when(site_full == "Phelps Road" ~ "low",

```

```

        site_full == "North of Phelps Bridge" ~ "low",
        site_full == "South of Venoco Bridge" ~ "moderate",
        site_full == "South of Phelps Bridge" ~ "moderate",
        site_full == "Devereux Creek" ~ "moderate",
        site_full == "North of Venoco Bridge" ~ "high",
        site_full == "Main Channel" ~ "high")) |>
mutate(taxon = str_replace_all(taxon, "_", " "), taxon = str_to_title(taxon)) |>
filter(taxon %in% focus_taxa) |>
# create column that calculates log-transformed abundance
mutate(log_ave_lit = log(ave_lit + 1)) |>
mutate(bins = factor(bins, levels = c("low", "moderate", "high")))


# base layer: ggplot
ins_bins_plot <- ggplot(data = aq_ins_clean_bins,
                        mapping = aes(x = bins, # x-axis: DO bins
                                      y = log_ave_lit, # y-axis: log_ave_lit
                                      color = taxon)) + # color by taxon

# first layer: points show individual abundance measurements
geom_point(shape = 16, # shape is circle
           size = 2, # size is 2
           alpha = 0.7) + # opacity is 0.7
# calculate median for each bin for each taxon
stat_summary(fun = median,
            geom = "point",
            color = "red",
            size = 3) +
# separate panels for each taxon
facet_wrap(~ taxon) +
# use non-default colors for points
scale_color_viridis_d() +
labs(x = "DO levels", # relabel x-axis
     y = "log + 1 transformed average abundance/L", # relabel y-axis
     # creates descriptive title
     title = "Invertebrate abundance vs DO levels") +
# uses non-default theme
theme_bw() +
# remove redundant legend
theme(legend.position = "none")

# shows plot
ins_bins_plot

```

Invertebrate abundance vs DO levels



```
# Kruskal-Wallis test between abundance and bins for Copepods
kruskal.test(log_ave_lit ~ bins,
             data = aq_ins_clean_bins,
             subset = taxon == "Copepod") # Kruskal-Wallis chi-squared = 4.405, df = 2, p-value = 0.1105
```

Kruskal-Wallis rank sum test

data: log_ave_lit by bins

Kruskal-Wallis chi-squared = 4.405, df = 2, p-value = 0.1105

no significant difference

```
# Kruskal-Wallis test between abundance and bins for Ostracods
kruskal.test(log_ave_lit ~ bins,
             data = aq_ins_clean_bins,
             subset = taxon == "Ostracod") # Kruskal-Wallis chi-squared = 6.7235, df = 2, p-value = 0.034
```

Kruskal-Wallis rank sum test

data: log_ave_lit by bins

Kruskal-Wallis chi-squared = 6.7235, df = 2, p-value = 0.03467

```
# significant difference of abundance btwn DO levels

# Pairwise test between abundance and individual bins for Ostracods
bins_ostracods <- aq_ins_clean_bins |>
  filter(taxon == "Ostracod")
pairwise.wilcox.test(
  bins_ostracods$log_ave_lit,
  bins_ostracods$bins,
  p.adjust.method = "bonferroni") # p = 0.051 for comparison btwn moderate and high DO levels
```

Pairwise comparisons using Wilcoxon rank sum test with continuity correction

data: bins_ostracods\$log_ave_lit and bins_ostracods\$bins

	low	moderate
moderate	0.387	-
high	1.000	0.051

P value adjustment method: bonferroni

From this figure and the medians, it looks like Copepods and Ostracods prefer moderate levels of DO. The others don't appear to have a preference, but nematode abundance is especially low throughout so there might just not be enough to tell for that taxon. Only Ostracod abundance medians for each DO level bin significantly differ according to a Kruskal-Wallis test. A pairwise comparison using Wilcoxon rank sum test found that abundance compared between moderate and high levels of DO for Ostracods has a p-value of 0.051. Otherwise, pairwise comparisons are not significant ($p > 0.1$).


```
# Made a copy of the composition plot but changed the bins to match the 4 quartiles of the data
composition_plot_1 <- DO_taxa |>
  mutate(
    do_category = case_when(
      med_do <= quantile(med_do, 0.25, na.rm = TRUE) ~ "Low DO",
      med_do <= quantile(med_do, 0.75, na.rm = TRUE) ~ "Medium DO",
      TRUE ~ "High DO"
    ),
    do_category = factor(do_category,
```

```

                                levels = c("Low DO",
                                             "Medium DO",
                                             "High DO"))
) |>
group_by(do_category, taxon) |>
summarize(
  total_abundance = sum(ave_lit, na.rm = TRUE),
  .groups = "drop"
) |>
ggplot(aes(x = do_category,
           y = total_abundance,
           fill = taxon)) +
geom_col(position = "fill") +
scale_y_continuous(labels = scales::percent_format()) +
labs(
  x = "Dissolved oxygen category",
  y = "Relative community composition",
  fill = "Taxon",
  title = "Invertebrate community composition shifts across dissolved oxygen levels"
) +
theme_classic()

# display plot
composition_plot_1

```



```
# running chi-square tests to compare proportional abundance across DO levels
copepods_matrix <- matrix(c(160.133333, 1112.133333, 2155.200000, 949.866667, 17.466667, 165.066667),
  dimnames = list(Oxygen = c("Low", "Moderate", "High"),
    Taxa = c("Copepod", "Other")))
view(copepods_matrix)
chisq.test(copepods_matrix)
```

Pearson's Chi-squared test

```
data: copepods_matrix
X-squared = 1297.8, df = 2, p-value < 2.2e-16
```

```
# X-squared = 1297.8, df = 2, p-value < 2.2e-16
# Proportion of Copepod abundance changes across DO levels.
pairwise_copepods <- pairwise.prop.test(copepods_matrix, p.adjust.method = "bonferroni")
pairwise_copepods
```

Pairwise comparisons using Pairwise comparison of proportions

```
data: copepods_matrix
```


	Low	Moderate
Moderate	<2e-16	-
High	0.89	<2e-16

P value adjustment method: bonferroni

```
# Significant difference between Low & High, Moderate & High

ostracods_matrix <- matrix(c(592.666667, 679.599963, 211.466667, 2893.600033, 153.733333, 28.800000),
  dimnames = list(Oxygen = c("Low", "Moderate", "High"),
    Taxa = c("Copepod", "Other")))
view(ostracods_matrix)
chisq.test(ostracods_matrix)
```

Pearson's Chi-squared test

```
data: ostracods_matrix
X-squared = 1318.3, df = 2, p-value < 2.2e-16
```

```
# X-squared = 1318.3, df = 2, p-value < 2.2e-16
# Proportion of Ostracod abundance changes across DO levels.
pairwise_ostracods <- pairwise.prop.test(ostracods_matrix, p.adjust.method = "bonferroni")
pairwise_ostracods
```

Pairwise comparisons using Pairwise comparison of proportions

```
data: ostracods_matrix
```

	Low	Moderate
Moderate	<2e-16	-
High	<2e-16	<2e-16

P value adjustment method: bonferroni

```
# Significant difference between all pairs of DO levels.

cladocera_matrix <- matrix(c(509.333333, 762.933333, 688.400000, 2417.066667, 10.000000, 155.066667),
  dimnames = list(Oxygen = c("Low", "Moderate", "High"),
```

```
Taxa = c("Copepod", "Other"))
view(cladocera_matrix)
chisq.test(cladocera_matrix)
```

Pearson's Chi-squared test

```
data: cladocera_matrix
X-squared = 184.58, df = 2, p-value < 2.2e-16
```

```
# X-squared = 184.58, df = 2, p-value < 2.2e-16
# Proportion of Cladocera abundance changes across DO levels.
pairwise_cladocera <- pairwise.prop.test(cladocera_matrix, p.adjust.method = "bonferroni")
pairwise_cladocera
```

Pairwise comparisons using Pairwise comparison of proportions

```
data: cladocera_matrix
```

	Low	Moderate
Moderate	< 2e-16	-
High	< 2e-16	4.2e-06

P value adjustment method: bonferroni

```
# Significant difference between all pairs of DO levels.
boatman_matrix <- matrix(c(10.133333,1262.133333, 50.000000,3055.466667, 1.333333,181.2), nrow = 3,
  dimnames = list(Oxygen = c("Low", "Moderate", "High"),
    Taxa = c("Copepod", "Other")))
view(boatman_matrix)
chisq.test(boatman_matrix)
```

Pearson's Chi-squared test

```
data: boatman_matrix
X-squared = 5.0378, df = 2, p-value = 0.08055
```

```
# X-squared = 5.0378, df = 2, p-value = 0.08055
# Proportion of Boatman abundance does not change across DO levels.
```

When comparing medians of taxon abundance across DO bins, only Ostracod abundance is significantly different between the moderate and high DO bins. When comparing proportions of taxon abundance across DO bins, the proportions of Copepods, Ostracods, and Cladocera significantly differ across the different DO bins. For Copepods, the differences exist between Low & High and Moderate & High DO levels. For Ostracods and Cladocera, the differences exist between all pairings of DO levels.

Chi-square

```
# create abundance table for chi-square test
chi_square_table <- aq_ins_clean_bins |>
  filter(taxon %in% focus_taxa) |>
  group_by(bins, taxon) |>
  summarise(
    total_abundance = sum(ave_lit, na.rm = TRUE),
    .groups = "drop"
  ) |>
  pivot_wider(
    names_from = taxon,
    values_from = total_abundance,
    values_fill = 0
  ) |>
  column_to_rownames("bins") |>
  as.matrix()

# show table used in chi-square test
chi_square_table
```

	Cladocera	Copepod	Hemiptera	Corixidae	Boatman	Nematode	Ostracod
low	980.2667	115.8667			0.2666667	0.1333333	1202.933333
moderate	510.6667	1170.4000			83.8666667	0.2666667	1057.866667
high	264.2667	1453.4667			8.8000000	0.0000000	7.866667

```
# chi-square test for differences in taxon composition across DO bins
chi_square_test <- chisq.test(chi_square_table)
```

```
# show chi-square result  
chi_square_test
```

Pearson's Chi-squared test

```
data:  chi_square_table  
X-squared = 2859.2, df = 8, p-value < 2.2e-16
```

The chi-square test suggests that taxon composition differs strongly across DO bins: the relative abundance of Cladocera, Copepod, Boatman, Nematode, and Ostracod is not evenly distributed across low, moderate, and high DO groups, $X^2 = 2859.2$, $df = 8$, $p < 2.2e-16$